

1. Pitfalls, False and Hidden Assumptions in Distributed Systems

Distributed systems often rely on several assumptions that may not hold in real-world deployments:

- Reliable Network: Developers often assume the network will never drop or delay packets, but in reality, networks are prone to failures and congestion.
- Zero Latency: There is a hidden assumption that communication between nodes is instantaneous, which neglects real-world delays.
- Infinite Bandwidth: Assumes no limitation on data transmission speed, which is unrealistic especially in large-scale systems.
- Always Available Nodes: Assumes that all nodes are always up, failing to consider system crashes or power failures.
- Global Clock: Distributed systems often presume perfectly synchronized clocks, which is practically impossible without introducing clock drift.
- Identical Performance: Assumes all nodes perform equally well, ignoring hardware and network discrepancies.
- Security: Assumes internal network communication is safe, potentially overlooking malicious attacks.

2. Distributed vs Decentralized Systems

- Distributed System: Comprises multiple independent components located on different machines that communicate and coordinate their actions to appear as a single coherent system. A central authority may still exist.

Example: Google Cloud

- Decentralized System: A subset of distributed systems where control is distributed; no single entity has full control.

Example: Blockchain, BitTorrent

Feature	Distributed System	Decentralized System
Control	May have central control	No central authority
Coordination	Coordinated actions	Autonomous peer actions
Failure Impact	Partial or complete	Localized

3. Mobile Cloud Computing vs Mobile Edge Computing

- Mobile Cloud Computing (MCC) involves offloading data processing and storage to centralized cloud servers. While it supports scalability and flexibility, it often suffers from high latency and dependence on connectivity.

- Mobile Edge Computing (MEC) brings computation and storage closer to the user, often at the edge of the network (e.g., base stations). This improves responsiveness and supports real-time applications.

Feature	Mobile Cloud	Mobile Edge
Latency	Higher	Lower
Processing Site	Remote cloud servers	Nearby edge devices
Use Cases	Cloud backup, Email	AR/VR, Self-driving cars

4. Grid Computing

Grid computing is a form of distributed computing where resources from multiple domains are pooled together to reach a common objective. It allows tasks to be broken down and processed in parallel across multiple machines.

- Example: SETI@home
- Key Characteristics: Resource sharing, distributed nodes, heterogeneous systems.

5. Cluster Computing

Cluster computing involves a group of tightly coupled computers that work together to perform tasks as a single system. These systems are usually located close together and connected via high-speed networks.

- Example: Hadoop clusters, Beowulf clusters
- Key Characteristics: High availability, load balancing, parallel processing

6. Parallel Computing

Parallel computing is the simultaneous use of multiple compute resources to solve a computational problem:

- Types: Data Parallelism, Task Parallelism
- Goal: Increase speed and efficiency
- Example: Multi-core processors handling matrix multiplication

7. Policies vs Mechanisms

- Policies are the "what" rules that govern system behavior (e.g., which process to run next).
- Mechanisms are the "how" implementation that enforces policies (e.g., context switching).
- Example: Scheduling policy (round robin) vs scheduling mechanism (timer interrupt).

8. Scale in Distributed Systems

Scalability refers to the system's ability to handle growth:

- Size Scalability: Adding more nodes
- Geographical Scalability: Nodes spread across locations
- Administrative Scalability: Support multiple organizations
- Example: Cassandra's ring topology supports horizontal scaling.

9. Ubiquitous Systems

Also known as pervasive computing, it aims to make computing available everywhere and anytime through embedded sensors and smart devices.

- Examples: Smart homes, wearable health monitors, IoT devices
- Features: Context-awareness, invisibility, mobility

10. Mobile Computing

Mobile computing allows users to access computational resources without a fixed location using portable devices.

- Components: Mobile devices, wireless networks, cloud services
- Examples: Smartphones, laptops with 4G/5G access

11. System Availability Calculation

Given subsystem availabilities: 60%, 70%, 80%, 90%

Availability in parallel:

$$A = 1 - (1 - A_1)(1 - A_2)(1 - A_3)(1 - A_4)$$

$$A = 1 - (0.4 * 0.3 * 0.2 * 0.1) = 1 - 0.0024 = 0.9976 = 99.76\%$$

12. P2P Architectural Structures

- Pure P2P: All nodes are equal (BitTorrent)
- Hybrid P2P: Centralized indexing with P2P content sharing (Napster)

Diagrams:

- Pure P2P: Mesh network where every node connects to others
- Hybrid P2P: Clients connect to server for metadata, then peer-to-peer for files

13. Hybrid System Architectures

Combines client-server and P2P or edge-cloud structures:

- Example 1: Edge devices perform initial processing; cloud stores data
- Example 2: Spotify uses central server for discovery and P2P for streaming

14. Middleware Architectures

Middleware connects applications and services:

- Message-Oriented: Kafka, RabbitMQ
- Remote Procedure Call (RPC): gRPC, XML-RPC
- Object Middleware: CORBA, Java RMI
- Database Middleware: ODBC, JDBC

15. Project Relevance with Course Concepts

Your project on Cassandra touches many Distributed & Parallel concepts:

- Distributed Systems: Cassandras architecture spans multiple nodes
- Replication & Consistency: Testing eventual consistency models
- Scalability: Horizontal scaling of nodes
- CAP Theorem: Cassandra trades off consistency for availability (AP)
- Fault Tolerance: Node failure recovery via replication
- Parallel Querying: Multiple consistency-check operations run in parallel
- Middleware: Internal services coordinate read/write consistency across clusters